

Hiring Team,

I am writing to apply for the position related to GenAI. I am excited to apply because this role pairs reliable backend engineering with real GenAI experimentation, including agentic workflows and evaluation-aware systems (RAG/knowledge) that produce dependable outputs for customer relationship and sales/marketing use cases.

I am a Senior LLM / GenAI Engineer and Tech Lead with 5+ years in enterprise software and hands-on leadership of GenAI delivery. My recent work centers on production LLM systems, agentic workflows, and brownfield codebases, where AI must meet architecture, testing, and NFR bar—not demos alone.

In my current role I lead a team of five engineers building agentic and generative systems for a customer interactions and customer engagement systems codebase. The following are all LLM-first systems (prompting, structured outputs, tool use, and/or retrieval as needed—not rule-only automation):

1. Log analyzer — Uses LLMs to parse, summarize, and compare log streams; surface anomalies and recurring patterns; and suggest likely causes to speed triage and protect SLAs.
2. Single-click code generation from LLD — One action produces implementation code: LLMs read the LLD (plus config/constraints) and generate output aligned with in-house frameworks and architecture, so the build stays consistent with the documented design.
3. Chat-based component test generation (LLD-grounded) — A conversational flow where users refine tests in chat while the LLM stays anchored to the LLD and relevant code context—LLM-generated test scenarios that track the design instead of drifting from the spec.
4. Single-click microservice creation — LLMs produce service skeletons, configs, and pipeline definitions from inputs and templates in one flow—base code plus build/deploy pipeline wired to our standards.
5. LLD design assistant — Conversational LLM that guides authors through LLD sections, asks clarifying questions, and drafts or revises LLD content against our structure and vocabulary.
6. In-house interview management system — LLM-powered workflows for role-aligned questions, screening summaries, feedback synthesis, and candidate-facing text—an internal platform that uses models end-to-end in the interview lifecycle, not only static forms.

I own LLD/HLD reviews, planning, and code reviews with product and architecture. Alongside this work I have led NFR initiatives (e.g. single-transaction latency 18s → 2s using Dynatrace, AWR, and stack-level tuning) in large-scale production systems.

Would you be open to a conversation about how I'd build reliable, evaluation-aware GenAI agent workflows into your backend systems for customer interactions and sales/marketing use cases? Thank you for your consideration.

Sincerely,

Saurabh Maurya

8805639108 · saurabh11.maurya@gmail.com

LinkedIn: <https://www.linkedin.com/in/saurabhmaurya/>

GitHub: <https://github.com/saurabh-maurya>